

AI Engineer — Practical Assessment

Influencer / Boutique Candidate Scoring

7 calendar days · Deliverables: source code, README, one sample report, and a 5 minute screen recording.

The problem

Boutiquaat's flywheel runs on influencer-owned storefronts - GCC creators and celebrities who curate, market, and sell as commissioned Boutique owners, earning on attributed sales. Growth depends on continually finding the next Boutique owner.

We want to know which observable signals about a prospective Boutique owner - audience composition, content style, engagement quality, category and brand fit - actually predict onboarding success, and build something that scores and ranks candidates on that basis before Celebrity Management commits scarce onboarding capacity to them.

Your task

Build a **working Python prototype** that takes a candidate's public profile / content signals (an Instagram or TikTok handle, or a structured/mocked profile dataset) as input and outputs: a fit/success score, the key signals driving that score, and a concrete onboard / hold / pass recommendation. It must be demonstrable end to end on at least one real or representative candidate from raw profile data, through scoring, to a specific recommendation a Celebrity Management lead could act on tomorrow.

Everything else is your call - architecture, libraries, model choice (classical ML, LLM-assisted feature extraction, hybrid, or something else), what data you can realistically source vs. mock, what you score, what you leave out. **We are assessing your judgement, not your feature count.** A small system that works, is tested, and that you can defend beats a large one that is half-finished; tell us in the README what you deliberately cut and why.

If you cannot access real historical onboarding outcomes or paid APIs, mock them and label the mock clearly - please do not spend your own money.

Ground rules

Seven days from receipt. Public libraries and APIs are fair game; identify what you used. AI coding assistants are allowed — but we will ask you to walk through your choices and the alternatives you rejected, so submit only what you can defend. No credentials in the repository; include a .env.example.

How we evaluate

The reasoning behind your scoring and validation methodology followed by Code quality and ML/AI implementation, architecture then testing, documentation, and how clearly you explain the work. Documented limitations count in your favour; overstated claims count against you.

Submit to [contact] by [deadline]. Questions are welcome at any point.